

Online and scalable kernel methods

I. Santamaría, S. Van Vaerenbergh

GTAS, Universidad de Cantabria

14 de febrero de 2022

Maître Universitario Oficial **Data Science**



Contents

Introduction

Large-Scale Kernel Methods

KLMS

Introduction

LMS

KLMS

Conclusions

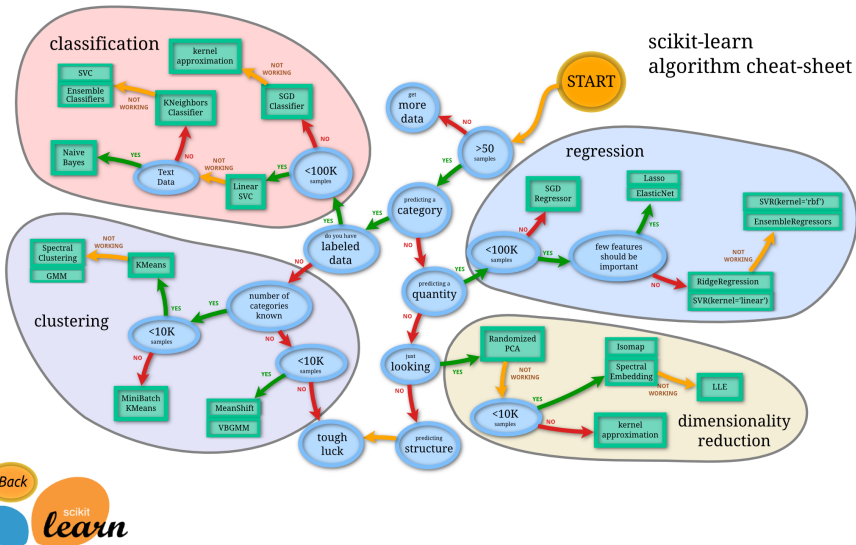
Introduction

- ▶ So far, we have studied **batch** kernel methods
- ▶ They need to store and compute an $n \times n$ kernel matrix \mathbf{K}

$$\mathbf{K} = \begin{bmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & k(\mathbf{x}_1, \mathbf{x}_2) & \cdots & k(\mathbf{x}_1, \mathbf{x}_n) \\ k(\mathbf{x}_2, \mathbf{x}_1) & k(\mathbf{x}_2, \mathbf{x}_2) & \cdots & k(\mathbf{x}_2, \mathbf{x}_n) \\ \vdots & \ddots & \ddots & \vdots \\ k(\mathbf{x}_n, \mathbf{x}_1) & k(\mathbf{x}_n, \mathbf{x}_2) & \cdots & k(\mathbf{x}_n, \mathbf{x}_n) \end{bmatrix}$$

- ▶ SVM, SVR: QP problem
- ▶ KRR, GP: inversion of \mathbf{K}
- ▶ Other kernel methods for clustering or dimensionality reduction compute the eigenvectors/eigenvalues of \mathbf{K}

High memory and computational requirements



- ▶ Problems with more than 100K patterns call for optimized implementations of kernel-based algorithms (e.g., LIBLINEAR for classification with linear SVMs)
- ▶ The number of support vectors grows linearly with the number of training patterns \Rightarrow Complexity of the trained machine

$$f(\mathbf{x}) = \sum_{i \in SVs} \alpha_i k(\mathbf{x}_i, \mathbf{x})$$

- ▶ How to make large-scale kernel machines?
 - ▶ Parallel computing: Multi-core, GPU
 - ▶ Approximate solutions:
 1. Random Fourier features
 2. Subsampling/sketching/chunking
 3. Low-rank approximations for \mathbf{K} , e.g., Nyström method

Linear vs. non-linear SVM

- ▶ For a linear SVM the optimal separating hyperplane can be expressed in closed form

$$\mathbf{w} = \sum_{i \in SVs} \alpha_i y_i \mathbf{x}_i \Rightarrow f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

- ▶ For a non-linear SVM, the optimal hyperplane in the feature space in general cannot be expressed in closed form
- ▶ We can only compute the output score using the kernel expansion

$$f(\mathbf{x}) = \sum_{i \in SVs} \alpha_i y_i k(\mathbf{x}_i, \mathbf{x})$$

- ▶ Training and testing are much faster with the linear SVM

Reprinted from: Yuan, G.-X., Ho, C.-H., Lin, C.-J. "Recent Advances of Large-Scale Linear Classification"(2012)

Data set	Linear	Non-linear (Gaussian kernel)
	Time / Accuracy	Time / Accuracy
MNIST38	0,1 / 96,82	38,1 / 99,70
ijcnn1	1,6 / 91,81	26,8 / 98,69
covtype	1,4 / 76,37	46.695,8 / 96,11
news20	1,1 / 96,95	383,2 / 96,90
real-sim	0,3 / 97,44	938,3 / 97,82
yahoo-japan	3,1 / 92,63	20.955,2 / 93,31
webspam	25,7 / 93,35	15.681,8 / 99,26

covtype: 581.012 patrones, 54 features

yahoo-japan: 176.103 patrones, 832.026 features

Random Fourier Features

- Standard kernel approach

$$\mathbf{x} \in \mathcal{R}^d \longrightarrow \Phi(\mathbf{x}) \quad \Phi(\mathbf{x})^T \Phi(\mathbf{y}) = k(\mathbf{x}, \mathbf{y})$$

- We could map the data explicitly to a low-dimensional feature space so that the inner product approximates the kernel function

$$\mathbf{x} \in \mathcal{R}^d \longrightarrow g(\mathbf{x}) \in \mathcal{R}^D \quad g(\mathbf{x})^T g(\mathbf{y}) \approx k(\mathbf{x}, \mathbf{y})$$

- How to choose a good mapping that approximates $k(\mathbf{x}, \mathbf{y})$?

Bochner Theorem

Any continuous shift-invariant kernel $k(\mathbf{x}, \mathbf{y}) = k(\mathbf{x} - \mathbf{y})$ is the Fourier transform of a probability density function $p(\omega)$

$$k(\mathbf{x}, \mathbf{y}) = \int p(\omega) e^{j\omega^T(\mathbf{x}-\mathbf{y})} d\omega = \mathbb{E} \left[e^{j\omega^T(\mathbf{x}-\mathbf{y})} \right]$$

- ▶ The Fourier transform of a Gaussian kernel is also a Gaussian function
- ▶ Since both the probability distribution and the kernel are real, the integral converges when the complex exponentials are replaced with cosines

- ▶ Draw D i.i.d. samples $\omega_1, \dots, \omega_D$ from $p(\omega)$ (a d -variate Gaussian pdf). Note that ω_i has the same dimension as \mathbf{x} .
- ▶ Draw D i.i.d. samples b_1, \dots, b_D from a uniform distribution on $[0, 2\pi]$

$$g(\mathbf{x}) = \sqrt{\frac{2}{D}} [\cos(\omega_1^T \mathbf{x} + b_1) \quad \dots \quad \cos(\omega_D^T \mathbf{x} + b_D)]^T$$

- ▶ Then, $g(\mathbf{x})^T g(\mathbf{y})$ is an estimator of $k(\mathbf{x}, \mathbf{y})$
- ▶ RFF provide an explicit mapping with some advantages
 1. We can work with a linear SVM/SVR
 2. If $D \ll n$ we can approximate the kernel matrix by the rank- D matrix $\mathbf{G}\mathbf{G}^T$, where \mathbf{G} has the D -dimensional RFFs as rows

Subsampling

- ▶ Typically the kernel expansion is sparse (only a few $\alpha_i \neq 0$)

$$f(\mathbf{x}) = \sum_{i \in SVs} \alpha_i k(\mathbf{x}_i, \mathbf{x})$$

- ▶ Choose a good subset of training patterns: by KNN or heuristic techniques
- ▶ Incremental training, mini-batches, chunking or sketching
 - ▶ Split the training data set into N mini-batches:
 $\mathcal{S} = \{\mathcal{S}_1, \dots, \mathcal{S}_N\}$
 - ▶ SVM trained with the first mini-batch $\{\mathcal{S}_1\} \rightarrow \# SV_1$
 - ▶ SVM trained with $\{\mathcal{S}_2, SV_1\} \rightarrow \# SV_2$
 - ▶ SVM trained with $\{\mathcal{S}_3, SV_2\} \rightarrow \# SV_3$
 - ▶ ...

Low-rank approximation of \mathbf{K}

- ▶ To improve scalability of kernel methods we can also approximate the rank- n kernel matrix \mathbf{K} by a low-rank approximation $\tilde{\mathbf{K}}$ with rank $r \ll n$
- ▶ Recall that RFFs provide a low-rank approximation $\tilde{\mathbf{K}} = \mathbf{G}\mathbf{G}^T$
- ▶ What is the advantage of working with $\tilde{\mathbf{K}}$ instead of \mathbf{K} ?
 - ▶ Memory requirements reduce from $n(n+1)/2$ to nr
 - ▶ For KRR we need to compute

$$\alpha = (\tilde{\mathbf{K}} + \lambda \mathbf{I})^{-1} \mathbf{y}$$

if $\tilde{\mathbf{K}}$ has rank $r \ll n$, its inverse can be computed more efficiently

Nyström method

- Pick $r \ll n$ columns at random from \mathbf{K}

$$\mathbf{K} = \begin{matrix} & \begin{matrix} \text{\textit{C}} \\ \text{\textit{C}} \end{matrix} \\ \begin{matrix} \text{\textit{K}} \\ \text{\textit{K}} \end{matrix} = & \begin{matrix} n \\ n \end{matrix} & \left[\begin{array}{cc} \boxed{\begin{matrix} \mathbf{K}_{11} \\ \mathbf{K}_{12} \end{matrix}} & \begin{matrix} \mathbf{K}_{21}^T \\ \mathbf{K}_{22} \end{matrix} \end{array} \right] \\ & \begin{matrix} \text{\textit{r}} \\ \text{\textit{r}} \end{matrix} \end{matrix}$$

- The low-rank Nyström approximation is

$$\tilde{\mathbf{K}} = \mathbf{C}\mathbf{K}_{11}^{-1}\mathbf{C}^T$$

Online learning

- ▶ We have studied **batch** kernel methods
 - ▶ Model parameters are obtained from a training data set and kept fixed during testing
 1. **Expansion coefficients**

$$\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_M)^T$$

2. **Dictionary** = Support Vectors

$$\mathcal{D} = \{\mathbf{x}_i, i = 1, \dots, M\}$$

- ▶ For a new test pattern \mathbf{x}_n the kernel method outputs

$$f(\mathbf{x}_n) = \sum_{i=1}^M \alpha_i k(\mathbf{x}_i, \mathbf{x}_n) = \mathbf{k}_n^T \boldsymbol{\alpha}$$

- ▶ In many learning problems patterns arrive sequentially:
online learning, **sample-by-sample learning**, **sequential learning**

- ▶ Both the expansion coefficients and the dictionary change with each new incoming pattern

1. Coefficients

$$\alpha(n) = (\alpha_1(n), \dots, \alpha_{M_n}(n))^T$$

2. Dictionary

$$\mathcal{D}_n = \{\mathbf{x}_i(n), i = 1, \dots, M_n\}$$

Note: the dictionary size M_n may change as well!

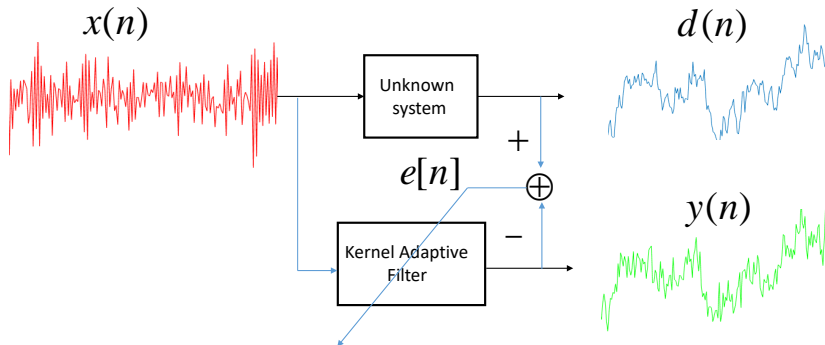
- ▶ The output is

$$f(\mathbf{x}_{n+1}) = \sum_{i=1}^{M_n} \alpha_i(n) k(\mathbf{x}_i(n), \mathbf{x}_{n+1})$$


- ▶ **Kernel Adaptive Filtering (KAF) / Online kernel methods:**
Updating algorithms for

$$\alpha_i(n) \rightarrow \alpha_i(n+1) \quad \mathcal{D}_n \rightarrow \mathcal{D}_{n+1}$$



Kernel Adaptive Filtering (KAF)



Time-embedding



Essentially, filtering is regression with time-embedding



Kernel Adaptive Filtering (KAF) problem

- ▶ With data observed up to $n - 1$ we build a kernel-based (GP, KRR, SVR) regressor or predictor
- ▶ Predictor parameters at $n - 1$: $\mathcal{D}_{n-1} = \{\mathbf{x}_i(n - 1)\}$, and $\alpha_i(n - 1)$
- ▶ We observe \mathbf{x}_n and make a new prediction

$$f(\mathbf{x}_n) = y(n) = \sum_i \alpha_i(n - 1) k(\mathbf{x}_i(n - 1), \mathbf{x}_n)$$


- ▶ Then, we get the desired output, $d(n)$, and compute the prediction error: $e(n) = d(n) - y(n)$
- ▶ Adaptive algorithm:

$$\alpha_i(n - 1) \rightarrow \alpha_i(n) \quad \mathbf{x}_i(n - 1) \rightarrow \mathbf{x}_i(n)$$


to minimize an error loss function (e.g., $\sum_n e(n)^2$)

Static vs. dynamic/adaptive models

Stationary processes → static models




Non-stationary processes → Dynamic/adaptive models



- ▶ With non-stationary time series or dynamic models, adaptive algorithms must be able to “forget” the past
- ▶ Ability to perform **tracking** under dynamic systems or to react to changes in the input signal properties
- ▶ Idea: past data are weighted less than more recent data


$$Cost = \sum_{k=0}^n \lambda^{n-k} e^2(k)$$

where $0 < \lambda < 1$ is an exponential forgetting factor



Least Mean Square (LMS) algorithm

- ▶ The most popular adaptive algorithm for linear systems is the **Least Mean Square (LMS)** algorithm (Widrow & Hopf, 1960)
- ▶ It is a **Stochastic Gradient Descent (SGD)** algorithm



Formulation

- ▶ At time n the vector of regressor coefficients (filter) is

$$\mathbf{w}(n) = [w_1(n) \quad w_2(n) \quad \dots \quad w_P(n)]^T$$

- ▶ The output is

$$y(n) = \mathbf{w}(n)^T \mathbf{x}_n = [w_1(n) \quad w_2(n) \quad \dots \quad w_P(n)] \begin{bmatrix} x(n) \\ x(n-1) \\ \dots \\ x(n-P+1) \end{bmatrix}$$

- ▶ In the linear case P is the dimension of the regressor and the time-embedding applied to the input
- ▶ We want to minimize

$$J(\mathbf{w}(n)) = \frac{1}{2} e(n)^2 = \frac{1}{2} \left(d(n) - \mathbf{w}(n)^T \mathbf{x}_n \right)^2$$

- The gradient is

$$\nabla J(\mathbf{w}(n)) = \begin{bmatrix} \frac{\partial J(\mathbf{w}(n))}{\partial w_1(n)} \\ \vdots \\ \frac{\partial J(\mathbf{w}(n))}{\partial w_P(n)} \end{bmatrix} = -2e(n) \begin{bmatrix} x(n) \\ \vdots \\ x(n-P+1) \end{bmatrix} = -e(n)\mathbf{x}_n$$

i.e., $\nabla J(\mathbf{w}) = -\text{error} \times \text{input vector}$

- We update the coefficients as

$$-\mu \nabla J(\mathbf{w}) = \mu e(n)\mathbf{x}_n$$

where μ is the

- Step-size (signal processing)
- Learning rate (machine learning)

Least Mean Square (LMS) Algorithm

Initialize $\mathbf{w}(0) = [0, \dots, 0]^T$

For $n = 0, \dots$


$$e(n) = d(n) - \mathbf{w}(n)^T \mathbf{x}_n$$

$$\mathbf{w}(n+1) = \mathbf{w}(n) + \mu e(n) \mathbf{x}_n$$


end

- ▶ Advantages:
 - ▶ It works with non-stationary signals (tracking)
 - ▶ Easy implementation and very low computational cost:
 $P + 1$ multiplications and $P - 1$ additions per iteration
- ▶ Drawbacks:
 - ▶ Slow convergence


Example



1 simulación



200 simulaciones promediadas
(Monte Carlo)



Kernel LMS

- ▶ Can we apply the LMS algorithm in the feature space?
- ▶ The answer is yes and the resulting algorithm is the **Kernel Least Mean Square (KLMS)**
- ▶ We apply a linear adaptive filter in the transformed (feature) space
- ▶ In the input space we have a non-linear adaptive filter

- ▶ The input patterns are mapped to a high-dimensionality space

$$\mathbf{x}_n \rightarrow \Phi(\mathbf{x}_n)$$

- ▶ Let us denote the linear regressor or filter in the feature space as $\omega(n)$
- ▶ The output (prediction) is

$$y(n) = \omega(n)^T \Phi(\mathbf{x}_n)$$

- ▶ After observing the desired output, $d(n)$, we compute the error, $e(n) = d(n) - \omega(n)^T \Phi(\mathbf{x}_n)$, and update the filter coefficients with the LMS

$$\omega(n+1) = \omega(n) + \mu e(n) \Phi(\mathbf{x}_n)$$

- ▶ Initialize $\omega(0) = \mathbf{0}$
- ▶ Applying the LMS sequentially we get

$$\begin{aligned}\omega(n) &= \omega(n-1) + \mu e(n-1) \Phi(\mathbf{x}_{n-1}) \\ &= [\omega(n-2) + \mu e(n-2) \Phi(\mathbf{x}_{n-2})] + \mu e(n-1) \Phi(\mathbf{x}_{n-1}) \\ &= \omega(n-2) + \mu [e(n-2) \Phi(\mathbf{x}_{n-2}) + e(n-1) \Phi(\mathbf{x}_{n-1})] \\ &= \omega(0) + \mu \sum_{k=0}^{n-1} e(k) \Phi(\mathbf{x}_k) = \mu \sum_{k=0}^{n-1} e(k) \Phi(\mathbf{x}_k)\end{aligned}$$

- ▶ The output can be expressed through the **kernel trick**

$$\begin{aligned}y(n) &= \omega(n)^T \Phi(\mathbf{x}_n) = \sum_{k=0}^{n-1} \mu e(k) \Phi(\mathbf{x}_k)^T \Phi(\mathbf{x}_n) = \\ &= \sum_{k=0}^{n-1} \underbrace{\mu e(k)}_{\alpha(k)} k(\mathbf{x}_k, \mathbf{x}_n)\end{aligned}$$

- ▶ Each new input pattern is added to the dictionary

$$\mathcal{D}_n = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$$

- ▶ The coefficient associated to each pattern is $\alpha(n) = \mu e(n)$
- ▶ The output is the kernel expansion

$$y(n+1) = \sum_{k=1}^n \alpha(k) k(\mathbf{x}_k, \mathbf{x}_{n+1})$$

- ▶ **Problem:** The dictionary size and the complexity of the expansion grow unbounded with n

Limiting the dictionary growth

- ▶ In practice, KLMS algorithms apply different methods to limit the dictionary size
 - ▶ Sliding window \rightarrow fixed budget
 - ▶ A new pattern is added to the dictionary only when some criterion is met
 - ▶ A popular criterion is the **coherence**

Coherence

Given a dictionary $\mathcal{D}_M = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_M\}$, and a kernel normalized to fulfill $k(\mathbf{x}, \mathbf{x}) = 1$, coherence is defined as

$$\eta = \max_{i \neq j} |k(\mathbf{x}_i, \mathbf{x}_j)|, \quad \forall (\mathbf{x}_i, \mathbf{x}_j) \in \mathcal{D}_M$$

Coherence-based KLMS

- ▶ Given a size M dictionary: \mathcal{D}_M
- ▶ For each new input pattern \mathbf{x}_n
 1. If $\max_{i \in \mathcal{D}_M} |k(\mathbf{x}_i, \mathbf{x}_n)| \geq \eta_0$
 - ▶ Update expansion coefficients

$$\alpha(n+1) = \alpha(n) + \mu e(n) \mathbf{k}(\mathcal{D}_M, \mathbf{x}_n)$$

2. If $\max_{i \in \mathcal{D}_M} |k(\mathbf{x}_i, \mathbf{x}_n)| < \eta_0$
 - ▶ Add \mathbf{x}_n to the dictionary



$$\mathcal{D}_{M+1} = \{\mathcal{D}_M, \mathbf{x}_n\}$$

- ▶ Update expansion coefficients

$$\alpha(n+1) = \begin{bmatrix} \alpha(n) \\ 0 \end{bmatrix} + \mu e(n) \mathbf{k}(\mathcal{D}_{M+1}, \mathbf{x}_n)$$



Example

Prediction of the Mackey-Glass chaotic time-series




KLMS vs LMS

Inversion of a non-linear system (non-linear equalization)





KIN40K dataset

- Forward kinematics of an eight-link all-revolute robot arm



- Data set: 40000 data
 - Input: 8 features (angles)
 - Output: one-dimensional
- Non-linear regression problem

Comparison KLMS vs. coherence-based KLMS



Conclusions

- ▶ We can improve kernel methods scalability by
 - ▶ Random Fourier features
 - ▶ Subsampling/sketching
 - ▶ Low-rank approximations of the kernel matrix
- ▶ Online kernel methods
 - ▶ Sample-by-sample (sequential) adaptation of the model parameters: dictionary+coefficients
 - ▶ A popular algorithm is the KLMS: dictionary growth
 - ▶ KLMS + coherence criterion: limits the dictionary size
 - ▶ Many other KAF algorithms:
<https://github.com/steven2358/kafbox> (toolbox en Matlab)