

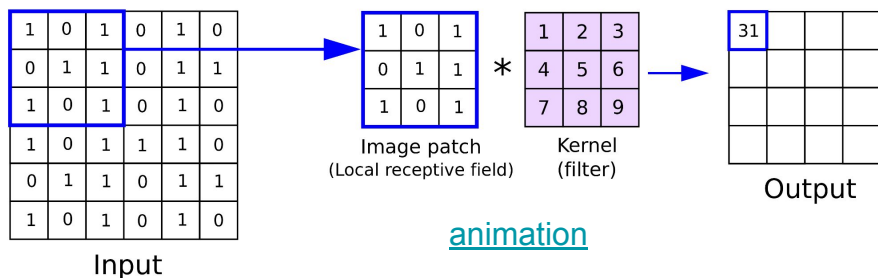
Image classification

Master InterUniversitario de Data Science
Santander, Spain
March 2022

Ignacio Heredia
iheredia@ifca.unican.es
Instituto de Física de Cantabria
(CSIC-UC)


CNNs theory - Convolutional Layers

The convolution operation




[animation](#)

Example. Edge detection



In addition to the kernel size you can set the **stride** (jump step) and the **padding** (fill border with zeros).

Animations: [\[1\]](#), [\[2\]](#), [\[3\]](#)




We have to repeat the operation across the depth of the feature map.

[animation](#)

CNNs theory - Convolutional Layers

Inside the same convolutional layer we can have *multiple kernels*. Each kernel is specialized in one task and is learnt through backpropagation. Each kernel will produce a different feature map.




The convolution operation is *spatially invariant*, making it well suited for images (ie. we want to detect edges no matter where they are located).

Additionally, reusing parameters (one simple 3×3 kernel instead of full dense layer) reduces the networks size and avoids overfitting.

Convolutional Neural Networks have been nevertheless successfully used with other types of data (eg. time series).

CNNs theory - Pooling Layers

One Feature Map




(a) Max-Pooling (stride 2)

Pooling layers reduce the spatial dimensions of the feature maps (the outputs of the convolutional layers).

They can be omitted if one uses conv layers with large stride.

One Feature Map




(b) Average-Pooling (stride 2)


CNNs theory - Overall structure

The idea is to use alternatively **convolutional** layers and **pooling** layers (optional) until we reach the situation of having *many feature maps* with *low spatial size*. You can stack several conv layers (two 3×3 is better than one 5×5) before each pooling.

We flatten the last feature maps and apply some **dense** layers till we reach the output. Then apply softmax if we are doing classification.



CNNs theory - Overall structure


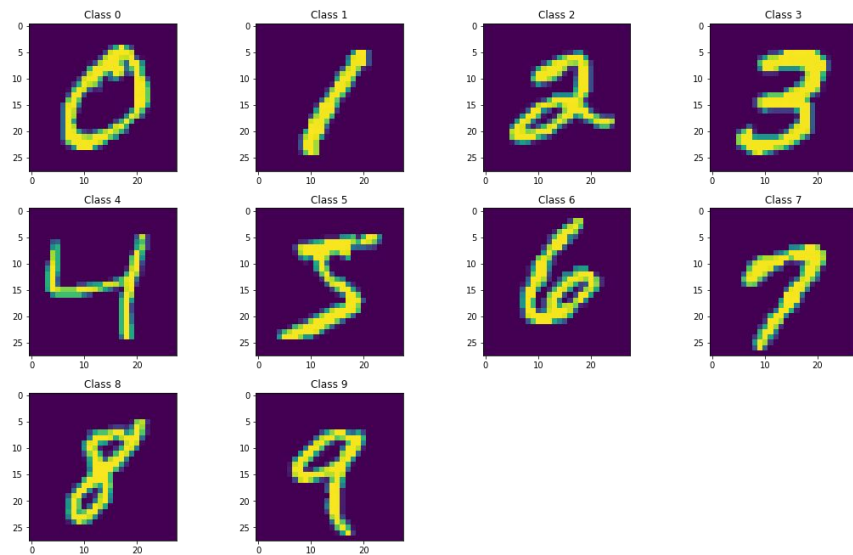


eg. detect simple edges

Progressively, as one goes deeper into the CNN, feature maps should encode *more complex/abstract* information.

eg. detect cat tail

Exercise 1 - Visualization



Exercise 2 - Create a CNN




Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Exercise 3.1 - Test the functions

Regularizers

- 0) None
Loss: 2.2, Acc: 0.511
- 1) l1_l2
Loss: 1.5e+02, Acc: 0.126
- 2) l2
Loss: 6.6, Acc: 0.126
- 3) l1
Loss: 1.3e+02, Acc: 0.099

Initializers

- 0) he_uniform
Loss: 2.2, Acc: 0.364
- 1) RandomNormal
Loss: 2.2, Acc: 0.275
- 2) he_normal
Loss: 2.2, Acc: 0.231
- 3) TruncatedNormal
Loss: 2.3, Acc: 0.213
- 4) glorot_uniform
Loss: 2.3, Acc: 0.211
- 5) lecun_uniform
Loss: 2.3, Acc: 0.198
- 6) RandomUniform
Loss: 2.3, Acc: 0.18
- 7) VarianceScaling
Loss: 2.3, Acc: 0.174
- 8) glorot_normal
Loss: 2.4, Acc: 0.116
- 9) lecun_normal
Loss: 2.3, Acc: 0.11
- 10) Ones
Loss: 1.4e+01, Acc: 0.11
- 11) Zeros
Loss: 2.3, Acc: 0.099


Optimizers

- 0) Adagrad
Loss: 0.25, Acc: 0.926
- 1) Nadam
Loss: 0.26, Acc: 0.921
- 2) Adamax
Loss: 0.36, Acc: 0.889
- 3) RMSprop
Loss: 0.41, Acc: 0.88
- 4) Adam
Loss: 0.41, Acc: 0.871
- 5) Adadelata
Loss: 1.2, Acc: 0.601
- 6) SGD
Loss: 2.2, Acc: 0.371

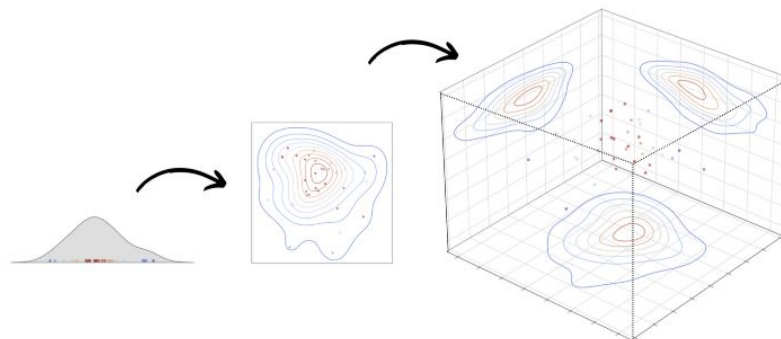
Exercise 3.2 - Hyperparameter search

Possible strategies

- Grid search ★
- Random search ★★
- Bayesian guided search ★★★




Bergstra, J., & Bengio, Y. (2012). [Random search for hyper-parameter optimization](#). Journal of machine learning research, 13(2).



Search space grows with the increased dimensionality of permissible hyperparameters ([ref](#))

Exercise 3.2 - Hyperparameter search

Grid search



Random search

